

Real-Time Programming Languages

Ada: Features for Real-Time and Safety

Martin Becker

TU München
Institute for Real-Time Computer Systems (RCS)

December 15, 2015



Lehrstuhl für
Realzeit-Computersysteme

Technische Universität München





- 1 Recap: Processes and Scheduling
 - Process vs. Thread
 - Scheduling
- 2 Ada's Real-Time & Safety Features
 - Ada Features
 - Exceptions
 - Real-Time Annex
 - Ravenscar Profile
- 3 Conclusion
- 4 References



Section 1

Recap: Processes and Scheduling

Process vs. Thread (1)

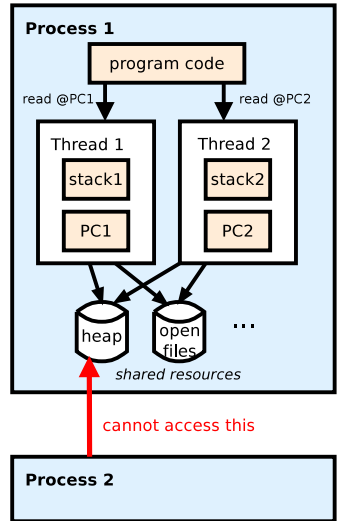
General Concept

Process

- one instance of your program code
- has its own memory space and resources (heap, open files, ...)
- consists of at least one *thread*
- other processes cannot access its resources

Thread

- a flow of execution \Rightarrow has its own program counter (PC) and stack
- runs in parallel to other threads



Remark: *Inter-Process Communication* for resource sharing between processes.

Process vs. Thread (2)



Implementation in Linux

- in Linux kernel **everything is a process**
- threads are *light-weight processes*
 - = a process sharing resources with other light-weight process(es)
- every process (light-weight or not), has its own PCB:
 - stack, program counter, ...
- processes can have children
 - can be light-weight or normal processes
 - in Linux, the “root” process is called `init`
 - everything starts here
 - try `ps tree -p 1`



Process vs. Thread (3)

Demo: creating a process (see `process.c` on Moodle)

Build+run `process-demo`:

```
1 $ gcc process.c o pdemo
$ ./pdemo

P: my process id is 14635
P: created child
6 P: created child
C: changed the variable to: 42
C: closed the file.
C: paused (press key)
C: exit
11 P: The variable is now: 9
P: Read from the file: t
P: exit
```

Inspect process tree:

```
2 $ pstree -p 14635
pdemo (14635) --- pdemo (14636)
```

Process has created a new child process.

This creates a new (heavy-weight) process, which is completely independent of the caller.



Process vs. Thread (4)

Demo: creating a thread (see `thread.c` on Moodle)

Build+run thread-demo:

```
$ gcc thread.c o tdemo
$ ./tdemo
3
P: my process id is 23375
P: created child
P: paused (press key)
C: my process id is 23376
8 C: changed the variable to: 42
C: closed the file.
C: paused (press key)
P: The variable is now: 42
P: READ failed: Bad file
    descriptor
```

Inspect process tree:

```
$ pstree -p 23375
tdemo(23375) --- tdemo(23376)
```

For Linux this is just another process.

This creates a new (light-weight) process, which shares resources with the caller.



Process vs. Thread (4)

Demo: creating a thread (see `thread.c` on Moodle)

Build+run thread-demo:

```
$ gcc thread.c o tdemo
$ ./tdemo
3
P: my process id is 23375
P: created child
P: paused (press key)
C: my process id is 23376
8 C: changed the variable to: 42
C: closed the file.
C: paused (press key)
P: The variable is now: 42
P: READ failed: Bad file
    descriptor
```

Inspect process tree:

```
$ pstree -p 23375
tdemo(23375) --- tdemo(23376)
```

For Linux this is just another process.

This creates a new (light-weight) process, which shares resources with the caller. **Remark:** The example uses a low-level API of Linux. Usually new threads are not created like this, but with the *pthread* library.




Summary:

In Linux everything is a process. We only need to deal with how to schedule processes.

Schedulability Analysis

- processes have to take turns on the CPU
- OS decides how they take turns, using a *scheduling policy*
 - round-robin, priority-based, earliest-deadline first, etc.
- in real-time systems, processes usually have *deadlines* - if they do not finish computing by then, something bad might happen
- **schedulability analysis** is a method to determine, whether all processes can keep their deadlines



$task := (P, e, D)$
 P = period (release time)
 e = WCET
 D = relative deadline



Context Switch

scheduling = switching between processes

- when suspending one process and waking up another, scheduler performs a *context switch*:

- 1 pause the currently running process
- 2 save the state of the retiring process to PCB:
 - PC, registers, I/O status, privileges, open files, ...
- 3 load state of the waking process from its PCB
- 4 resume the waking process

PCB in Linux (2% of it)

```
struct task_struct {  
    volatile long state;  
3 void *stack;  
    // ...  
    int on_cpu;  
    int prio;  
    pid_t pid;  
8 u64 start_time; /* monotonic  
    time in nsec */  
    /* open file information */  
    struct files_struct *files;  
    /* signal handlers */  
    struct signal_struct *signal;  
13 /* CPU-specific state of this  
    task */  
    struct thread_struct thread;  
};
```

(<https://github.com/torvalds/linux/blob/master/include/linux/sched.h>)



- switching takes time
 - several μs per switch, depending on data size (cache!)
 - several hundreds of switches per second (`grep ctxt /proc/<pid>/status`), depending on the scheduler settings
- switching between threads is much faster than switching between processes
- an incoming interrupt also causes a context switch



...one would talk about *mutex* and *semaphore*

- these are mechanisms provided by the OS
- they allow to synchronize threads
 - avoid race conditions on shared data
 - ...see Lecture 7



...one would talk about *mutex* and *semaphore*

- these are mechanisms provided by the OS
- they allow to synchronize threads
 - avoid race conditions on shared data
 - ...see Lecture 7

... but in Ada we have more high-level mechanisms which replace these

- internally, these may build on mutex and semaphore
- you may still use mutex and semaphore directly, if you insist (“Florist” library)



Section 2

Ada's Real-Time & Safety Features



- *certified* Ada compilers are available since 1983
- before Ada 95 the compilers had problems with tasking and efficiency
- since Ada 95, most innovations fall into the real-time area to fix these:
 - better concurrency model
 - new dispatching policies (such as non-preemptive, round-robin, EDF)
 - Timing Events
 - monitoring task execution times
 - high-integrity profiles, such as *Ravenscar*
- **today:** Ada is used for highly safety-critical systems

Famous users of Ada



Famous users:

- Air Traffic Management: DFS
- Aviation: Boeing 777 (99.9%!)
- Railway: TGV
- Rockets: Ariane
- Satellites: INMARSAT
- Banking: Reuters
- Medical: JEOL Nuclear MRI
- Military: Eurofighter combat aircraft
- . . .





User Reports:

- once it compiles, it works \Rightarrow less development time in total (more coding, much less debugging time)
- harder to mix different data types (units, casts, etc)
- programmers have to think harder, which yields better code

Moreover (today's lecture)

- amenable to static analysis \Rightarrow reduce testing effort
- small memory footprint
- Ada compilers usually have to pass a huge test suite (“ACVC”, $> 3,800$ tests), whereas almost all C compilers (including gcc) are known to have hundreds of bugs



In case of an unexpected program state

- 1 the control flow of the causing thread is interrupted
- 2 the control flow is handed over to an *exception handler*
- 3 if there is no handler, the thread terminates

Typical exceptions:

- `Constraint_Error`: exceeding variable's type bounds
- `Program_Error`:
- ...

Exceptions (2)



This causes an exception at line 5:

```
procedure main is
  type mydays is Integer range 1 .. 31
  d : mydays := 31;
begin
5   d := d + 1; -- raises an exception
   Put_Line("day is:" & d'img); -- never reached
end main;
```

Some exceptions are foreseen by the compiler. Others are not.
To handle them:

```
3   procedure main is
     type mydays is Integer range 1 .. 31
     d : mydays := 31;
     begin
       d := d + 1; -- raises an exception
       Put_Line("day is:" & d'img); -- never reached
     exception
8       -- control flow continues here
       -- do something to recover
     end main;
```



- some types of exceptions can be turned off by pragmas, e.g., Integer Overflow
- how are they implemented?
 - the compiler inserts checks into the program
 - if checks fail, program jumps to the exception handler
 - **therefore, exceptions consume memory space and execution time**
- in embedded systems: no space and no time
 - prove absence of exceptions
 - turn those off that can be proven
 - handle the remainder
 - Ada: `last_chance_handler`
 - one exception handler, global for the entire program
 - “last chance” before program terminates



```
with Ada.Real_Time;
```

- Ada implementations offering above package must comply to that annex (GNAT does, but coverage is platform-dependent)
- additional semantics, such as
 - integrated priority-based interrupt handling
 - deterministic scheduling via fixed and dynamic priority
 - preemptive scheduling and monitoring execution time
 - forbidden: `terminate`, `abort`, ...
 - protection against priority inversion (...Pathfinder Mars rover...)
 - minimizes risk of race conditions (...Therac-25 medical radiation machine...)
 - multi-processor support
 - ...



Interesting:

- `pragma Priority(...)`: Tasks can be given priorities
- `pragma Queuing_Policy(...)`: entry queues can support *priorities* on top of FIFO
- `pragma Max_Task_Entries(...)`: limit queue length
- *hi-res timer*: as described earlier (Clock, Time, Duration)
 - RM: “real time is defined to be the physical time as observed in the external environment” ⇒ **scheduled programming model**
 - implementation must document: upper bound on the size of a clock jump, rate drift, tick length, ...
- `Ada.Execution_Time.Timers`: get notified when task exceeds a certain execution time
- **clearly, if OS is used then it must be a Real-Time OS**



- Ada is a language with many features
- some features are hard to analyze (`select...`)
- the run-time libraries of Ada can incur high cost
- the language offers to ban features by putting *restriction pragmas*
 - a program that uses a restricted feature will not compile
 - unused features are removed from the run-time \Rightarrow smaller and faster program

To obtain efficient and analyzable programs for real-time systems, which features shall we turn off?

a village in Great Britain, location of 8th International Real-Time Ada Workshop (IRTAW), 1997.



a village in Great Britain, location of 8th International Real-Time Ada Workshop (IRTAW), 1997.



- discussion at workshop about *what needs to be done in Ada, to target high-integrity, efficient real-time systems*
- participants agreed on a **list of forbidden features** in the Ada language, that hinder analysis of such systems \Rightarrow The *Ravenscar Profile*



`pragma Profile (Ravenscar);`

- Ada can be complex (`select`, `terminate`, ...)
- this profile: constrain tasking features to an analyzable subset
 - forbids `select`, `ATC`, `requeue`, `dynamic priors` ...
 - **this is why we do not look at those features in detail**
- **only concerned with tasking features of Ada**, nothing else
- goals:
 - full determinism
 - easier schedulability analysis
 - memory boundedness

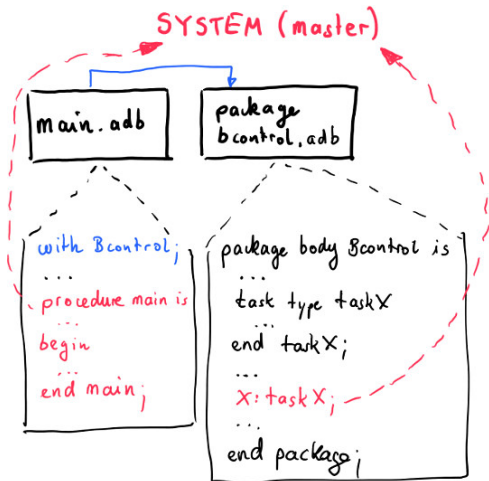


Tasks are either

- 1 *time-triggered* (periodic) or
 - released by a *single* delay until
- 2 *event-triggered* (sporadic)
 - released by a *single* protected entry
 - e.g., to realize interrupts

In addition:

- **no task hierarchies:** master of all tasks is “system”
- **no rendezvous:** all task interaction via *protected types*
- **static task set:** no dynamic task allocation, with static properties (e.g., task priorities), never terminate
- **entry queues:** have a capacity of exactly one task



- tasks are “all at same level”
- master is always the system
- for details see *blinker* code from Moodle

Ravenscar - Tasks, single release point



```
5 with Ada.Real_Time; use Ada.Real_Time;
task body my_periodic_task is
    Next_Time : Time;
    Period : constant Time_Span := Milliseconds(100);
begin
    Next_Time := Ada.Real_Time.Clock + Period;
    loop
        delay until Next_Time; -- without "until" is illegal!
        -- Do work...
    10 Next_Time := Next_Time + Period;
    end loop;
exception -- handler goes here
end my_periodic_task;
```

```
2 task body my_sporadic_task is
begin
    loop
        Monitor.Wait_Event; -- protected entry
        -- Do work...
    end loop;
7 exception -- handler goes here
end my_sporadic_task;
```

Ravenscar - Tasks, single release point



```
with Ada.Real_Time; use Ada.Real_Time;
2 task body my_periodic_task is
    Next_Time : Time;
    Period : constant Time_Span := Milliseconds(100);
begin
    Next_Time := Ada.Real_Time.Clock + Period;
7    loop
        delay until Next_Time; -- without "until" is illegal!
        -- Do work...
        Next_Time := Next_Time + Period;
    end loop;
12 exception -- handler goes here
end my_periodic_task;
```

```
task body my_sporadic_task is
2 begin
    loop
        Monitor.Wait_Event; -- protected entry
        -- Do work...
    end loop;
7 exception -- handler goes here
end my_sporadic_task;
```



- in embedded systems: a timer, a UART message, input capture, ...
- interrupts always share data with some other process \Rightarrow **race conditions**
- Ravenscar: protected object to encapsulate shared data
- ISR is a protected procedure, cannot be called directly
- no race condition possible because of mutual exclusion



- in embedded systems: a timer, a UART message, input capture, ...
- interrupts always share data with some other process \Rightarrow **race conditions**
- Ravenscar: protected object to encapsulate shared data
- ISR is a protected procedure, cannot be called directly
- no race condition possible because of mutual exclusion

```
protected Receiver is
  entry Wait (Msg: access Message);
  pragma Interrupt_Priority (
    UART_Priority);
private
  5  procedure My_ISR;
  pragma Attach_Handler(My_ISR,
    UART_Data_Arrival);
  Buffer: Contents(Message_Size);
  Length: Natural := 1;
  Message_Ready : Boolean := False;
  10 end Receiver;
```

- pragma: ensure that ceiling priority of handler \geq max. possible prio of the interrupt to be handled



...and the body:

```
protected body Receiver is
  procedure My_Isr is -- called when UART gets data
  begin
    -- ... access hardware, set 'Length' and 'Buffer'
    Message_Ready := True;
  end My_Isr;

  entry Wait (Msg: access Message) when Message_Ready is
  begin
    -- copy data to caller
    Msg.Value (1 .. Length) := Buffer (1 .. Length);
    Msg.Length := Length;
    -- reset buffer, it's empty now:
    Length := 1;
    Message_Ready := False;
  end Wait;
end Receiver;
```



... and the matching main task:

```
task Message_Processor is
  pragma Priority (System.Priority'First);
3 end Message_Processor;
task body Message_Processor is
  Next_Message : aliased Message (Size => Message_Size'Last)
  ;
begin
  loop
8   Receiver.Wait (Next_Message'Access); -- "pointer"
    -- do something in response...
  end loop;
end Message_Processor;
```

Note the *single release point* is the entry function in this example. This is a *sporadic* task, since we do not know when it will execute.

Ravenscar - Scheduling Policy for Tasks



Ravenscar: FIFO_Within_Priorities, preemptive



Additionally: Static CPU assignment on multi-core systems.

⇒ OS must support this. Linux: kernel RT patchset



Ravenscar: Ceiling_Locking

- =Priority Ceiling Protocol (see Lecture 7)
- protected objects get a ceiling priority, which is higher than that of all tasks that are accessing it
- caller inherits the priority of the protected object



- highest-priority runnable task in the system can always run
- no mutex locks or semaphores required
- tasks never block when calling a protected object
- can be analyzed with simple algorithms such as *rate-monotonic* schedulability test:

$$\text{task set schedulable if } u = \sum_{i=1}^n \frac{e_i}{P_i} \leq n(\sqrt[n]{2} - 1) \quad (1)$$

u =processor load, n =number of tasks, $\forall i : D_i = P_i$



- memory for each task (e.g., stack) must be resolved latest at link time
- i.e., tasks and protected types must not be dynamically allocated
- *implicit* memory allocation not allowed (compiler must obey this)
- **no restrictions** for large or dynamic-size objects, i.e., keyword `new` is allowed



Queues in protected objects can hold at most one task

- i.e., each entry inside can only have one caller
- (procedures or functions do not have queues)
- program will compile, but ...
- if one task is waiting in the queue and another one wants to enqueue \Rightarrow Exception `Program_Error`

Protected objects cannot have more than one entry

- multiple functions and procedures are okay
- \Rightarrow at most one task can be waiting for a guard

workarounds:

- see *Guide for the use of the Ada Ravenscar Profile in high-integrity systems*



pragma Detect_Blocking is part of Ravenscar:

- *potentially* blocking operations cannot be called from within protected objects
 - a lot of them are already banned by the profile: *select*, *task entry*, *delay*, ...
- however, reading from files, etc. could block
- currently, this is a run-time check
 - with this pragma, every potentially blocking operation raises an exception *when called*
 - compiler does not check this!
- **simplifies WCET computation**



- all communication/synchronization between tasks must be protected:
 - via protected object or suspension object or atomic object
 - only one task at a time can work on shared data
 - access is serialized through FIFOs
 - \Rightarrow no race conditions
- no potentially blocking operations from protected entries, and PCP
 - a task that is blocking other tasks cannot get blocked again
 - no “cycles” possible
 - \Rightarrow no deadlocks possible



- there are more forbidden features:

```
No_Dynamic_Attachment, No_Local_Protected_Objects,  
No_Local_Timing_Events, No_Requeue_Statements,  
No_Specific_Termination_Handlers, Simple_Barriers,  
Max_Entry_Queue_Length => 1, Max_Protected_Entries => 1,  
Max_Task_Entries => 0, No_Dependence =>  
Ada.Asynchronous_Task_Control, No_Dependence =>  
Ada.Calendar, No_Dependence =>  
Ada.Execution_Time.Group_Budgets, No_Dependence =>  
Ada.Execution_Time.Timers, No_Dependence =>  
Ada.Task_Attributes, No_Dependence =>  
System.Multiprocessors.Dispatching_Domains
```

- as Ada grows, new features enter this list



The Guide for the use of the Ada Ravenscar Profile in high integrity systems, A. Burns, B. Dobbing and T. Vardanega, 2003 provides a rationale for each of the restrictions.



With the described task model, scheduling policy and other restrictions, implementation is easy:

- little or no RTOS required
- minimal program is only $\approx 2\text{kB}$ of object code
- deterministic tasking
- reduced footprint of run-time
 - smaller run-time means lower certification cost
 - means smaller memory and thus less power consumption
 - means shorter boot time
- such implementations are well-studied and **certifiable up to highest assurance levels**
 - e.g., “DO-178B Level A” for commercial avionics (failure of level-A software results in catastrophic consequences, such as aircraft crash or unintended release of weapons)



- annex “Safety and Security”/“High Integrity Systems” provides *pragmas* that the developer can turn on:
 - `Normalize_Scalars`: set uninitialized variables to an out-of-range default value
 - `Reviewable`: provide implementation information for analysis, e.g.,
 - execution time, memory usage and mapping from source to object codes, presence of run-time checks, ... ⇒ **useful for WCET analysis**
 - many others (no protected types, no allocators (`new` etc), no exceptions, no float, no fixed, no delay, no recursion, detect blocking) ...
- attribute `'valid` to check whether value of object is within legal range, e.g., to recognize bit flips from cosmic rays, hardware errors etc.



Section 3

Conclusion



For safety-critical real-time systems in Ada, use

- a compiler supporting **annex Real-Time** \Rightarrow to get appropriate features
 - requires Real-Time OS
- the **Ravenscar profile** \Rightarrow ban tasking features to get determinism and enable schedulability analysis
 - enforces a simple program structure
 - requires Real-Time kernel (scheduling policy “FIFO within Priorities”) to work correctly
- a compiler supporting **annex High Integrity Systems** \Rightarrow to get information about the object code and restrict language even further, if desired



- verification of Ravenscar programs
 - static analysis of concurrent code
 - scheduling analysis
 - (WCET)
- dynamic and static analysis of sequential Ada code
 - sequential code is not addressed in Ravenscar
- formal verification (similar to Model Checking in Esterel) for Ravenscar programs
 - all-in-one for the restricted Ravenscar subset



Section 4

References



- *Ada95 Lovelace tutorial*, David A. Wheeler (who wrote that in his free time), <http://www.adahome.com/Tutorials/Lovelace/lovelace.htm>
- *Ada Wikibook*, http://en.wikibooks.org/wiki/Ada_Programming
- *The Boeing 777 Flies on 99.9% Ada*, <http://archive.adaic.com/projects/atwork/boeing.html>
- *Ada 95 Eliminates Race Conditions*, J.G.P. Barnes, Parallel and Distributed Real-Time Systems, 1995.
- *Programming Real-Time with Ada 2005*, P. Rogers, <http://www.embedded.com/design/prototyping-and-development/4025713/Programming-real-time-with-Ada-2005>, 2006.
- *Guide for the use of the Ada Ravenscar Profile in high-integrity systems*, A. Burns et. al., University of York, Technical Report YCS-2003-348, 2003.
- *Ada 95 Rationale: The Language - The Standard Libraries*, J. Barnes, Springer, 1995.

All online resources as of December 2015.